

Inference Latency vs. Classification Macro-F1 Trade-Off Across Evaluated Architectures

